

Lecture 8:

Advanced Binary Trees

Our Roadmap

◆ Priority Queue (binary heap)

- ◆ Min-heap insert / delete-min


◆ Binary Heaps in Dynamic Arrays

- ◆ $O(n)$ time to build min-heap

◆ Binary Search Tree (BST)

- ◆ BST operators
- ◆ Balanced BST (AVL-tree)

$\{L \text{ Root } R\}$
inorder traversal



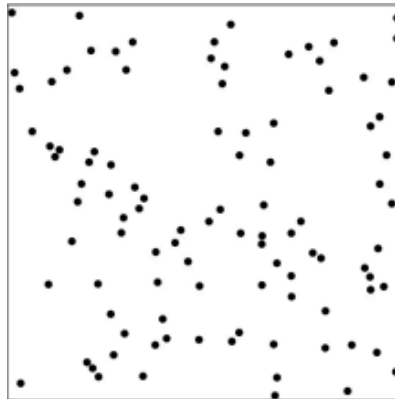
$a + b * c - d / e$

在树中 left 子树 在树中 right 子树

Priority Queue

- ◆ A priority queue stores a set S of n integers and supports the following operations:
 - ◆ *Insert(e)*: adds a new integer to S
 - ◆ *Delete-min*: removes the smallest integer in S , and returns it.
- ◆ Priority Queue applications:
 - ◆ Artificial intelligence (A* algorithm)
 - ◆ Operating systems (load balancing)
 - ◆ Graph searching (Shortest path algorithms)

8	4	7
1	5	6
3	2	



Priority Queue Example

- ◆ Suppose that the following integers are inserted into an initially empty priority queue
 - ◆ $S = \{93, 39, 1, 26, 8, 23, 79, 54\}$
 - ◆ Perform Delete-min, the operation returns 1, and $S = \{93, 39, 26, 8, 23, 79, 54\}$
 - ◆ Perform Delete-min, the operation returns 8, and $S = \{93, 39, 26, 23, 79, 54\}$
 - ◆ Perform
- ◆ Unlike an ordinary queue (FIFO), a priority queue guarantees that the elements always leave in ascending order (or descending order with *Delete-max*), regardless of the order by which they are inserted.

Priority Queue Implementation



- ◆ We will implement a priority queue using a data structure called the “binary heap” to achieve the following guarantees:
 - ◆ $O(n)$ space consumption
 - ◆ $O(\log n)$ insertion time
 - ◆ $O(\log n)$ delete-min time
- ◆ The binary heap data structure is an array object that we can view as a complete binary tree.
 - ◆ Level 0 to $h-1$ are full
 - ◆ Leaf nodes in level h are “as far left as possible”

Binary Heap 不唯一，顶是最小

- ◆ Let S be a set of n integers. A binary heap on S is a binary tree T satisfying:
 - ◆ (1) T is complete
 - ◆ (2) Every node u in T corresponds to a distinct integer in S , the integer is called the key of u (and is stored at u)
 - ◆ (3) If u is an internal node, the key of u is smaller than those of its child nodes
- ◆ Note that:
 - ◆ Condition 2 implies that T has n nodes
 - ◆ Condition 3 implies that the key of u is the smallest in the subtree of u

Binary Heap Example *delete min*

- Two possible binary heaps on $S = \{93, 39, 1, 26, 8, 23, 79, 54\}$:




- The binary heaps of a set S is not unique.
- The smallest integer of S must be the key of the root.

Binary Heap Insertion

- ◆ We perform $\text{insert}(e)$ on a binary heap T as follows:
 - ◆ Step 1: Create a leaf node z with key e , while ensuring that T is a complete binary tree, it means there is only one place where z could be added.
 - ◆ Step 2: Set $u \leftarrow z$
 - ◆ Step 3: If u is the root, return.
 - ◆ Step 4: If the key of $u >$ the key of its parent p , return
 - ◆ Step 5: Otherwise, swap the keys of u and p . Set $u \leftarrow p$, and repeat from Step 3.


Binary Heap Insertion

- Suppose we want to insert 15 into the binary heap below:



Binary Heap Insertion


- First add 15 as a new leaf, making sure that we still have a complete binary tree.



- Step 3 is not true, go to Step 4,
- Step 4 is not true, go to step 5.

Binary Heap Insertion


- First add 15 as a new leaf, making sure that we still have a complete binary tree.



- Swap the keys of u and its parent p
- Set $u \leftarrow p$, go back to Step 3)
- Step 3 and Step 4 are not true, go to Step 5.

Binary Heap Insertion

- First add 15 as a new leaf, making sure that we still have a complete binary tree.




- Swap the keys of u and p
- Set $u \leftarrow p$, go back to Step 3)
- Step 3 is not true, Step 4 is true, return. Insertion complete.

Binary Heap Delete-min

- ◆ We perform delete-min on a binary heap T as follows:
 - ◆ Step 1: Report the key of the root
 - ◆ Step 2: Identify the rightmost leaf z at the bottom level of T
 - ◆ Step 3: Delete z , and store the key of z at the root
 - ◆ Step 4: Set $u \leftarrow$ the root
 - ◆ Step 5: If u is leaf, return
 - ◆ Step 6: If the key of $u <$ the keys of the children of u , return
 - ◆ Step 7: Otherwise, let v be the child of u with a smaller key
Swap the keys of u and v . Set $u \leftarrow v$, and repeat from Step 5

Binary Heap Delete-min


- Assume that we perform a delete-min from the binary heap below:



- Delete-min delete root node, and we should maintain the rest nodes as a complete binary tree.

Binary Heap Delete-min


- ◆ First, find the rightmost leaf at the bottom level, it is node with key 79.



- ◆ Note that the tree is still a complete binary tree after removing this leaf.

Binary Heap Delete-min

- Remove the leaf, but place the key value 79 in the root.




min-heap: ① Insert(e) ~~Delete min~~

max-heap: ① Insert
② max-heap

- Step 4: set u as the root.
- Step 5 and 6 are not true,
- Go to Step 7.

Binary Heap Delete-min

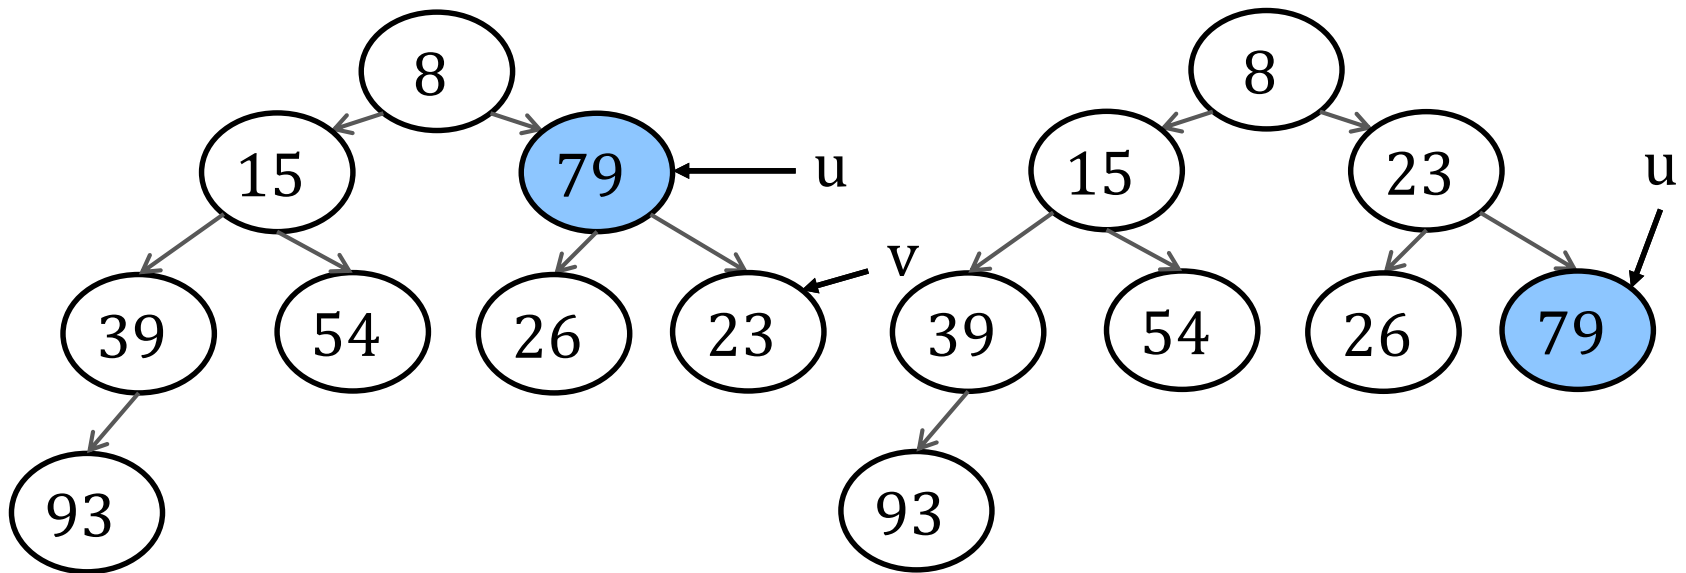
- Let v be the child of u with a smaller key.
- Swap the keys of u and v , and set $u \leftarrow v$



- Go to Step 5
- Step 5 and Step 6 are not true, go to Step 7

Binary Heap Delete-min


- ◆ Let v be the child of u with a smaller key.
- ◆ Swap the keys of u and v , and set $u \leftarrow v$



- ◆ Go to Step 5
- ◆ Step 5 is true, return. Delete-min complete.

How to find rightmost leaf?

- ◆ Before we analyzing the time complexity of insert and delete-min, let us first consider a sub-problem:
- ◆ Given a complete binary tree T with n nodes, how to identify quickly the rightmost leaf node at the bottom level of T (i.e., colored node in below tree).
 - ◆ It is Step 1 in insert algorithm, and Step 2 in delete-min algorithm



How to find rightmost leaf?

- ◆ We give a clever algorithm for solving the subproblem in $O(\log n)$ time.
- ◆ Write the value n in binary form. We can do that in $O(\log n)$ time.
- ◆ Skip the most significant bit. We will scan the remaining bits from left to right, start from root,
 - ◆ If the bit is 0, we go to the left child of the current node
 - ◆ Otherwise, go to right child

Find Rightmost Leaf Example

- ◆ Here $n = 9$, binary form: 1001

- ◆ Skip the first bit '1'


- ◆ We scan the remaining bits

- ◆ Start from root node 1.

- ◆ The 2nd leftmost bit is 0, so we visit left, and go to node 15

- ◆ The 3rd leftmost bit is 0, so we visit left, and go to node 39

- ◆ The 4th leftmost bit is 1, so we turn right, and go to node 79 (done).



Time Complexity Analysis

- ◆ We are now ready to prove that our insertion and delete-min algorithms finish in $O(\log n)$ time.
- ◆ It suffices to point out the key facts:
 - ◆ Step 1 of the insertion algorithm (page 8) and Step 2 of the delete-min algorithm (page 13) can be performed in $O(\log n)$ time, using our solution to previous sub-problem
 - ◆ The rest of insertion ascends a root-to-leaf path, while that of delete-min descends a root-to-leaf path. The time is $O(\log n)$ in both cases.
- ◆ Thus, we guarantee: (1) $O(n)$ space consumption, (2) $O(\log n)$ insertion / delete-min operations.

Our Roadmap

- ◆ Priority Queue (binary heap)
 - ◆ Min-heap insert / delete-min
- ◆ Binary Heaps in Dynamic Arrays
 - ◆ $O(n)$ time to build min-heap
- ◆ Binary Search Tree (BST)
 - ◆ BST operators
 - ◆ Balanced BST (AVL-tree)

Binary Heaps in Dynamic Arrays


- ◆ We have already learned that the binary heap serves as an efficient implementation of a priority queue. Our previous discussion was based on pointers (for getting a parent node connected with its children). In this lecture, we will see a “pointerless” way to implement a binary heap, which in practice achieves much lower space consumption
- ◆ We will also see a way to build a heap from n integers in just $O(n)$ time, improving the obvious $O(n \log n)$ bound.

Recall

- ◆ A **priority queue** stores a set S of n integers and supports the following operations:
 - ◆ *Insert(e)*: adds a new integer to S
 - ◆ *Delete-min*: removes the smallest integer in S , and returns it.
- ◆ Let S be a set of n integers. A **binary heap** on S is a binary tree T satisfying:
 - ◆ (1) T is complete
 - ◆ (2) Every node u in T corresponds to a distinct integer in S , the integer is called the key of u (and is stored at u)
 - ◆ (3) If u is an internal node, the key of u is smaller than those of its child nodes

Storing a Complete Binary Tree

- ◆ Storing a complete binary tree using an array
- ◆ Let T be any complete binary tree with n nodes, let us linearize the nodes in the following manner:
 - ◆ Put nodes at a higher level before those at a lower level
 - ◆ Within the same level, order the nodes from left to right
- ◆ Let us store the linearized sequence of nodes in an array A of length n . Example:



nik delete
 $n \log n$

$O(\log n)$ 时间
最右的 leaf

1	39	8	79	54	26	23	93
---	----	---	----	----	----	----	----

Storing in an array

Property 1

- ◆ Let us refer to the i -th element of A as $A[i]$, for simplicity, we assume the index of A starts from 1.
- ◆ Lemma: Suppose that node u of T is stored at $A[i]$. Then, the left child of u is stored at $A[2i]$, and the right child at $A[2i+1]$.
- ◆ Observe this from the example of the previous slide
- ◆ Proof leaves as your homework.
- ◆ Hints, consider the number of nodes after u , but before its left child.

More Properties

- ◆ The following is an immediate corollary of the previous lemma:
- ◆ Corollary: Suppose that node u of T is stored at $A[i]$. Then, the parent of u is stored at $A[\lfloor i/2 \rfloor]$.
- ◆ The following is a simple yet useful fact:
- ◆ Lemma: the rightmost leaf node at the bottom level is stored at $A[n]$.
- ◆ Now we have got everything we need to implement the insertion and delete-min algorithms on the array representation of a binary heap.

Insert 15

1	39	8	79	54	26	23	93
---	----	---	----	----	----	----	----

1	39	8	79	54	26	23	93	15
---	----	---	----	----	----	----	----	----

1	39	8	15	54	26	23	93	79
---	----	---	----	----	----	----	----	----

1	15	8	39	54	26	23	93	79
---	----	---	----	----	----	----	----	----

Delete-min

1	15	8	39	54	26	23	93	79
---	----	---	----	----	----	----	----	----

79	15	8	39	54	26	23	93
----	----	---	----	----	----	----	----

8	15	79	39	54	26	23	93
---	----	----	----	----	----	----	----

8	15	23	39	54	26	79	93
---	----	----	----	----	----	----	----

Performance Guarantees

- ◆ Combining our analysis on (i) binary heaps and (ii) dynamic arrays, we obtain the following guarantees on binary heap implemented with a dynamic array:
 - ◆ Space consumption $O(n)$
 - ◆ Insertion: $O(\log n)$ time amortized
 - ◆ Delete-min: $O(\log n)$ time amortized

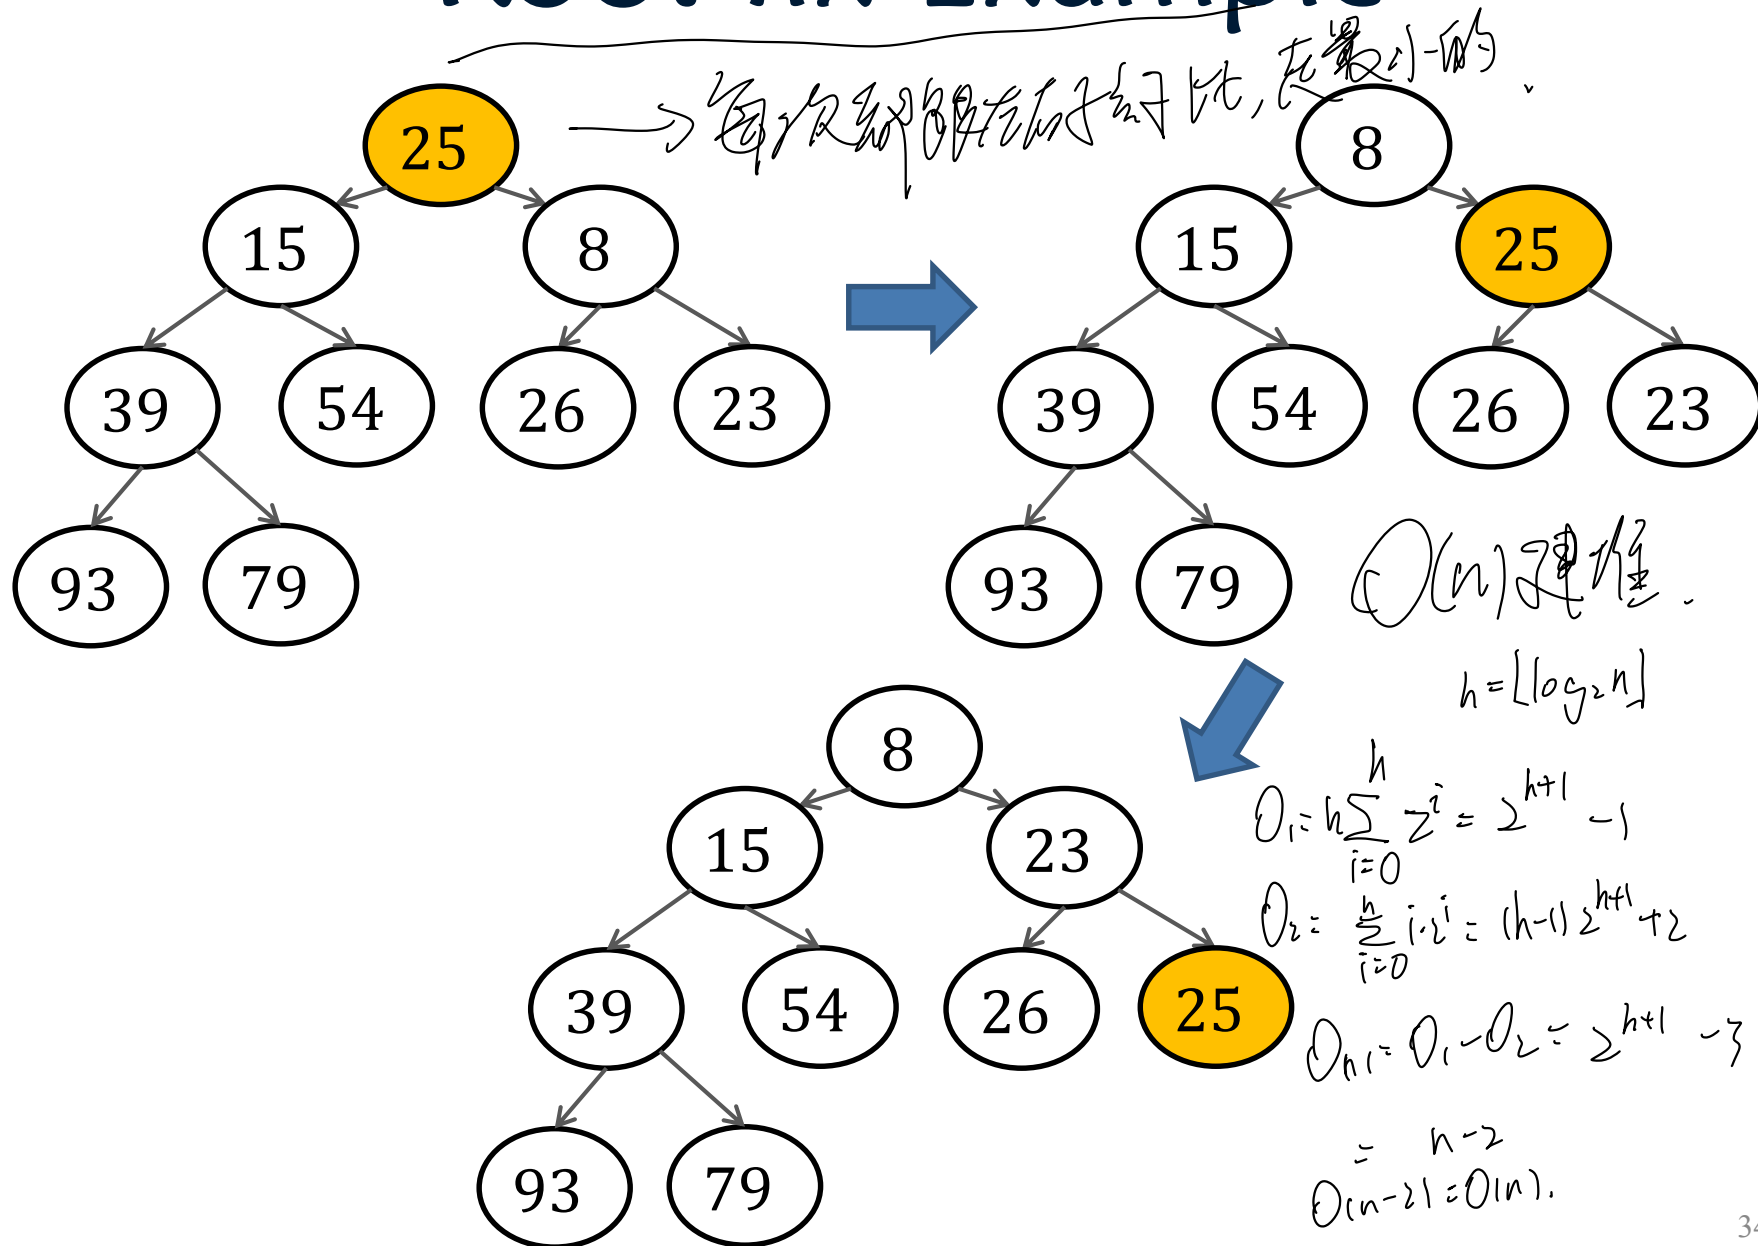
Build a binary heap in array

- ◆ Next we consider the problem of creating a binary heap on a set S of n integers. Obviously, we can do so in $O(n \log n)$ time by doing n insertions. However, this is an overall kill because the binary heap does not need to support any delete-min operations until all the n numbers have been inserted. This raises the question whether we can build the heap faster?
- ◆ The answer is positive: we will see an algorithm that does so in $O(n)$ time.

Root-fix operator

- ◆ We are given a complete binary tree T with root r . It guaranteed that:
 - ◆ The left subtree of r is binary heap
 - ◆ The right subtree of r is a binary heap
 - ◆ However, the key of r may not be smaller than the keys of its children.
- ◆ We define the root-fix operation, it fixes the issue and makes T a binary heap.
- ◆ Root-fix can be done in $O(\log n)$ time – in the same manner as the delete-min algorithm (step 4 - 7)

Root-fix Example



Building a Heap

- ◆ Create an array A that stores a set S of n integers, we can turn A into a binary heap on S using the following simple algorithm, which view A as a complete binary tree T :
- ◆ For each $i=n$ downto 1 :
 - ◆ Perform root-fix on the subtree of T rooted at $A[i]$
- ◆ Think: why are the conditions of root-fix always satisfied?

Building a Heap example

[illegible]

Root-fix

54	26	15	39	8	1	23	93
----	----	----	----	---	---	----	----

54	26	1	39	8	15	23	93
----	----	---	----	---	----	----	----

56	8	1	39	26	15	23	93
----	---	---	----	----	----	----	----

1	8	15	39	26	54	23	93
---	---	----	----	----	----	----	----

Complexity Analysis

- ◆ Lemma: The time complexity of turn array A into a binary heap on S is $O(n)$.

linear.

- ◆ Proof as follows:

- ◆ view A as a complete binary tree
- ◆ The height of T is h.
- ◆ Without loss of generality, assume that all the levels of T are full, i.e., $n=2^{h+1}-1$.
 - ◆ Why no generality is lost?
- ◆ Analyze the total running time of Build heap algorithm
- ◆ Proof that $\sum_{i=0}^h O(i * 2^{h-i}) = O(n)$ with $n=2^{h+1}-1$.

Our Roadmap

- ◆ Priority Queue (binary heap)
 - ◆ Min-heap insert / delete-min
- ◆ Binary Heaps in Dynamic Arrays
 - ◆ $O(n)$ time to build min-heap
- ◆ Binary Search Tree (BST)
 - ◆ BST operators
 - ◆ Balanced BST (AVL-tree)

Binary Search Tree (BST)

Binary Search Tree (especially, balanced BST) is the most powerful data structure of this course. This is without a doubt one of the most important data structures in computer science.

In extreme case, BST is equivalent to a linked list, thus, we guarantee the operations performance of BST by study AVL-tree.

Dynamic Predecessor Search

- ◆ Let S be a set of integers. We want to store S in a data structure to support the following operations:
 - ◆ A predecessor query: give an integer q , find its predecessor in S , which is the largest integer in S that does not exceed q .
 - ◆ Insertion: adds a new integer to S
 - ◆ Deletion: removes an integer from S
- ◆ Suppose that $S = \{3, 10, 15, 20, 30, 40, 60, 73, 80\}$
 - ◆ The predecessor of 23 is 20
 - ◆ The predecessor of 15 is 15
 - ◆ The predecessor of 2 does not exist
- ◆ Note that a predecessor query is more general than a “dictionary look-up”. Why?

Binary Search Tree (BST)

- ◆ We will learn a version of the BST that guarantees:

- ◆ $O(n)$ space consumption

- ◆ $O(h)$ time per predecessor query (hence, also per dictionary lookup)

- ◆ $O(h)$ time per insertion

- ◆ $O(h)$ time per deletion

} $\Rightarrow O(\log n)$



- ◆ where $n = |S|$, h is the height of BST, Note that all the above complexities hold in the worst case.

Binary Search Tree (BST)

- ◆ A BST on a set S of n integers in a binary tree T satisfying all the following requirements:
 - ◆ T has n nodes
 - ◆ Each node u in T stores a distinct integer in S , which is called the key of u
 - ◆ For every internal u , it holds that:
 - ◆ The key of u is larger than all the keys in the left subtree of u .
 - ◆ The key of u is smaller than all the keys in the right subtree of u .

BST Example

- Two possible BSTs on $S = \{3, 10, 15, 20, 30, 40, 60, 73, 80\}$




Predecessor Query

- ◆ Suppose that we have created a BST T on a set S of n integers. A predecessor query with search value q can be answered by descending a single root-to-leaf path:
 - ◆ (1) Set $p \leftarrow -\infty$ (p will contain the final answer at the end)
 - ◆ (2) Set $u \leftarrow$ the root of T
 - ◆ (3) If $u = \text{nil}$, then return p
 - ◆ (4) If key of $u = q$, then set p to q , and return p
 - ◆ (5) If key of $u > q$, then set u to the left child (now $u = \text{nil}$ if there is no left child), and repeat from Step (3)
 - ◆ (6) Otherwise, set p to the key of u and u to the right child, and repeat from Step (3)

Predecessor Query Example


- Suppose that we want to find the predecessor of 35



- Set $p \leftarrow -\infty$, $u = \text{root } 40$
- (3) and (4) are not true, go to (5)
- Since $40 > 35$, the predecessor cannot be in the right subtree of 40, so we move to the left child of 40, now $u = \text{node } 15$.

Predecessor Query Example


- Suppose that we want to find the predecessor of 35



- (3), (4) and (5) are not true, go to (6)
- Since $15 < 35$, $p \leftarrow 15$, since this is the predecessor of 35 so far.
- The predecessor cannot be in the left subtree of 15, so we move u to the right child, now $u = \text{node } 30$.

Predecessor Query Example

- Suppose that we want to find the predecessor of 35



- (3), (4) and (5) are not true, go to (6)
- Since $30 < 35$, $p \leftarrow 30$, since this is the predecessor of 35 so far.
- The predecessor will be in the right subtree of 30, but 30 does not have a right child. So algorithm terminates here with $p = 30$ as the final answer.

Time complexity Analysis

- ◆ Obviously, we spend $O(1)$ time at each node visited. Since the height of BST is h , therefore the total query time is $O(h)$.

Successor Query

- ◆ The opposite of predecessors are successors.
- ◆ The successors of an integer q in S is the smallest integer in S that is no smaller than q .
- ◆ Suppose that $S = \{3, 10, 15, 20, 30, 40, 60, 73, 80\}$
 - ◆ The successor of 23 is 30
 - ◆ The successor of 15 is 15
 - ◆ The successor of 81 does not exist
- ◆ Given an integer q , a successor query returns the successor of q in S .
- ◆ By symmetry, we know from the earlier discussion (on predecessor queries) that a successor query can be answered using a BST in $O(h)$ time.

BST Insertion


- ◆ Suppose that we need to insert a new integer e . First create a new leaf z storing the key e . This can be done by descending a root-to-leaf path:
 - ◆ 1. Set $u \leftarrow$ the root of T
 - ◆ 2. If $e < \text{the key of } u$
 - ◆ 2.1 If u has a left child, then set u to the left child
 - ◆ 2.2 Otherwise, make z the left child of u , and done
 - ◆ 3. Otherwise:
 - ◆ 3.1 If u has a right child, then set u to the right child
 - ◆ 3.2 Otherwise, make z the right child of u , and done.
 - ◆ Repeat from Step 2.
- ◆ The total cost is proportional to the height of T , i.e., $O(h)$

BST Insertion Example

1. 2. 3. 4. 5. 6. 7. 8. 9. 10. 11. 12. 13. 14. 15. 16. 17. 18. 19. 20. 21. 22. 23. 24. 25. 26. 27. 28. 29. 30. 31. 32. 33. 34. 35. 36. 37. 38. 39. 40. 41. 42. 43. 44. 45. 46. 47. 48. 49. 50. 51. 52. 53. 54. 55. 56. 57. 58. 59. 60. 61. 62. 63. 64. 65. 66. 67. 68. 69. 70. 71. 72. 73. 74. 75. 76. 77. 78. 79. 80. 81. 82. 83. 84. 85. 86. 87. 88. 89. 90. 91. 92. 93. 94. 95. 96. 97. 98. 99. 100.

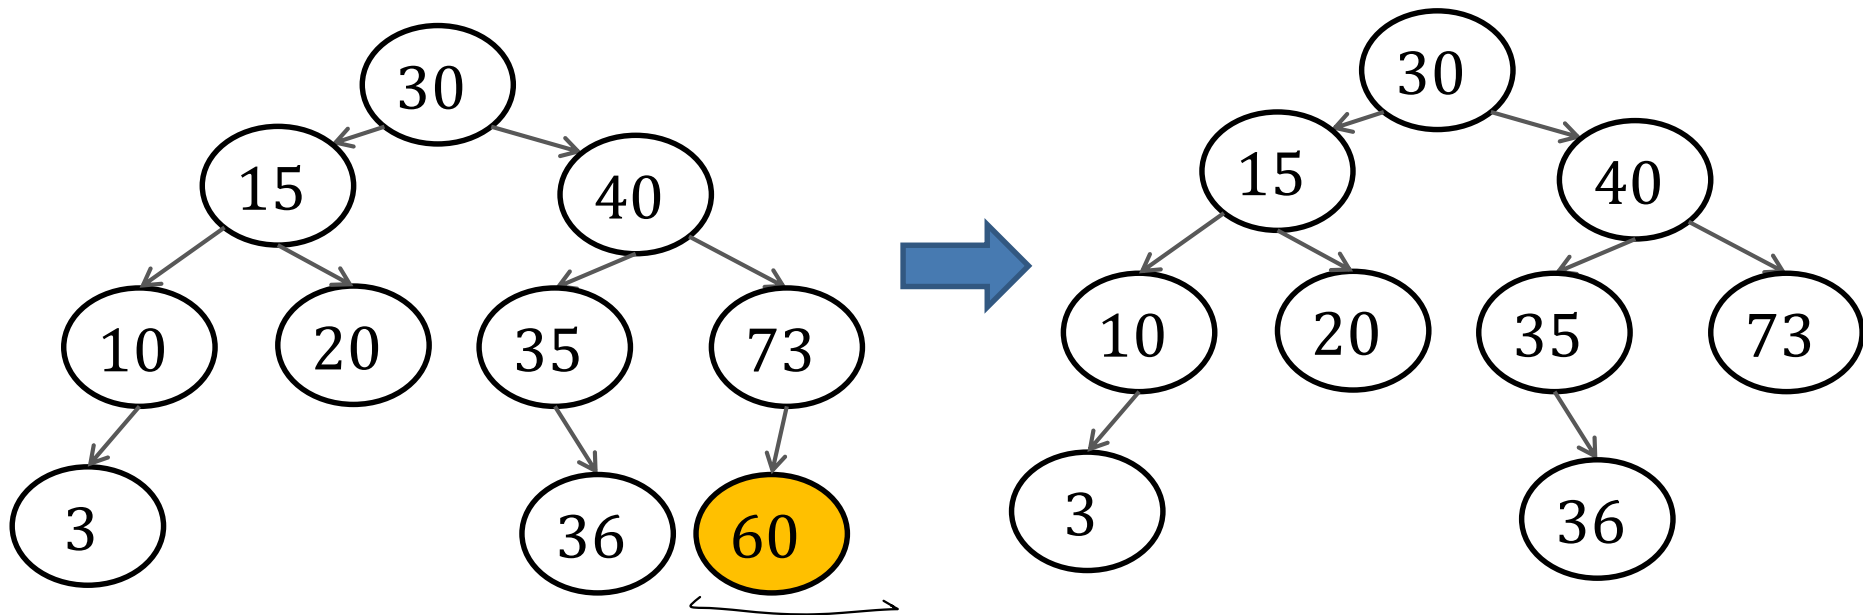
Case 1
Case 2 \rightarrow $\begin{cases} 2.1 \\ 2.2 \end{cases}$
Case 3

- Inserting 35:
- u is root 40, $e < \text{the key of u}$,
u has a left child, $u \leftarrow \text{node 15}$
- u is node 15, $e > \text{the key of u}$
u has a right child, $u \leftarrow \text{node 30}$
- u is node 30, $e > \text{the key of u}$,
u's right child is nil, then set z
as the right child of u. Done.



BST Deletion

- Suppose that we want to delete an integer e . First, find the node u whose key equals to e in $O(h)$ time (through a predecessor query).
- Case 1: if u is a leaf node, simply remove it from T .
- Example: remove 60



BST Deletion

- What happens if node u is not a leaf node?

- Case 2: if u has a right subtree:

- Find the node v storing the successor s of e .
- Set the key of u to s
- Case 2.1: if v is a leaf node, then remove it from T
- Case 2.2: otherwise, it must hold that v has a right child w , but not left child. Replace node v by subtree which rooted at w .


- Case 3: if u has no right subtree:

- It must hold that u has a left child v , Replace node u by the subtree rooted at v .

Case 2.1 Example *bst*

- ◆ Delete 40:
- ◆ u has a right subtree, node v (60) is the successor of 40.
- ◆ Set the key of u to 60
- ◆ v is a leaf, remove node v, done.

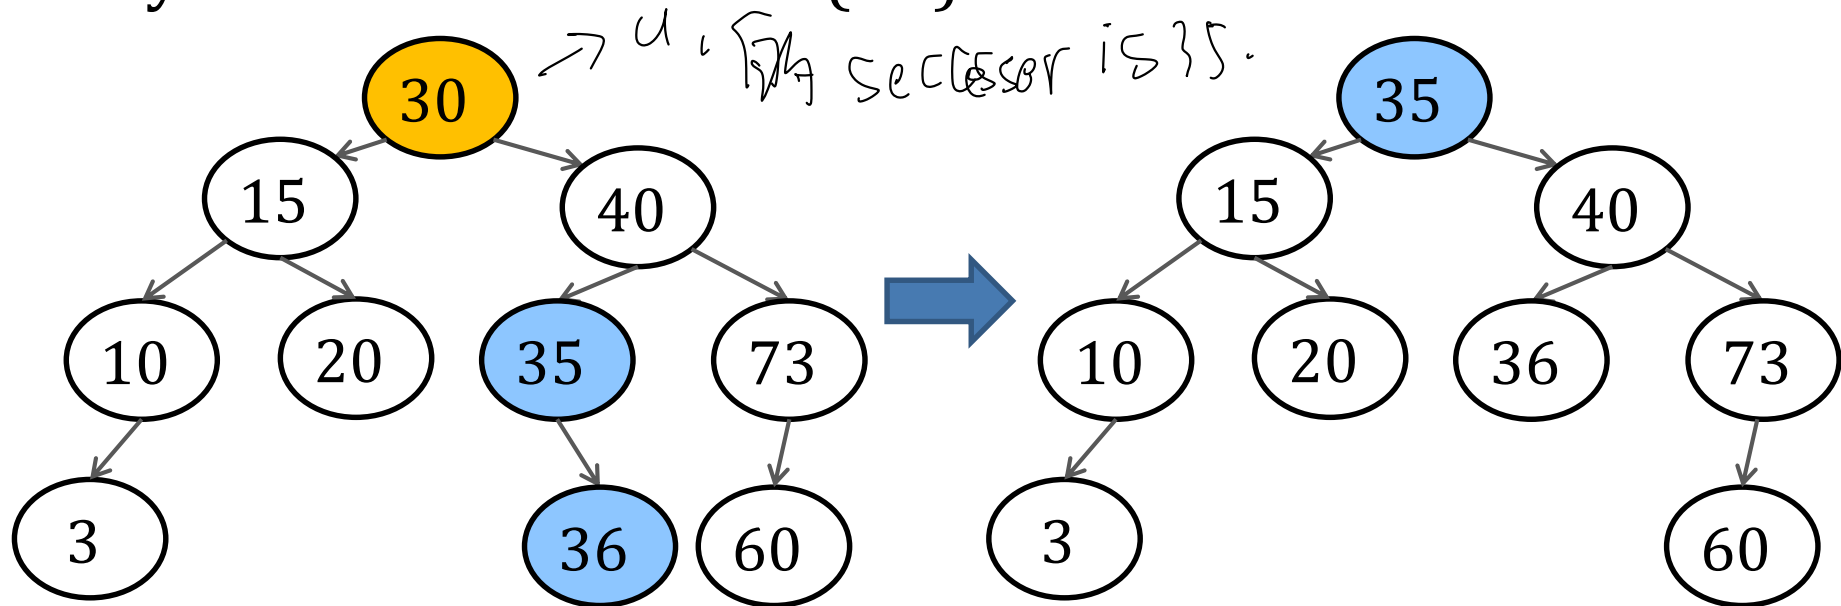
Case 2.1
1. Find the node
2. Set the key of u to v's



Case 2.2 Example


- ◆ Delete 30:
- ◆ u has a right subtree, node v (35) is the successor of 30.
- ◆ Set the key of u to 35
- ◆ v is not leaf node, it has right child w (36), replace node v by subtree rooted at w(36).

该节点不是叶子。



Case 3 Example

- ◆ Delete 73:
- ◆ u has no right subtree, and u must have a left child v (60), replace node u by node v(60).
- ◆ done.



BST Deletion


- ◆ In all above cases, we have essentially descended a root-to-leaf path (call it deletion path), and removed a leaf node.
- ◆ The cost so far is $O(h)$, recall that the successor of an integer can be found in $O(h)$ time.
- ◆ Given a set S of n integers, what is the maximum possible height of its BST?
 - ◆ $h = n$, why?
 - ◆ So what is the worst-case query cost? $O(n)$
 - ◆ However, we can guarantee $h = O(\log n)$ if the BST is balanced BST.

What is the height of tree


- ◆ Given a set S of n integers, what is the maximum possible height of its BST?
 - ◆ $h = n$, why?
- ◆ What is the worst-case query / insertion / deletion cost?
 - ◆ $O(n)$!!!
- ◆ How to achieve $O(\log n)$ time per operation?
 - ◆ Balanced Binary Search Tree

Balanced Binary Tree

- ◆ A binary tree T is balanced if the following holds on every internal node u of T :
 - ◆ The height of the left subtree of u differs from that the right subtree of u by at most 1.
- ◆ If u violates the above requirement, we say that u is imbalanced.



Balanced




Height of a Balanced Binary Tree

- ◆ Theorem: a balanced binary tree with n nodes has height $O(\log n)$.
- ◆ Proof. (left as homework)
- ◆ Hints:
 - ◆ 1) consider minimum number of nodes in a balanced binary tree with height h
 - ◆ 2) recursive equation
 - ◆ 3) analysis two cases: case 1) h is even, case 2) h is odd.
- ◆ With the height of balanced binary tree is $O(\log n)$, we can conclude that the cost of query operation is $O(\log n)$ on a balanced binary search tree.
- ◆ How about the cost of insertion and deletion on it?

Balanced BST

- ◆ An AVL-tree on a set S of n integers is a balanced binary search tree T , where the following hold on every internal node u
 - ◆ u stores the heights of its left and right subtrees.
- ◆ An AVL-tree on $S = \{3, 10, 15, 20, 30, 40, 60, 73, 80\}$
- ◆ For example, the number 3 and 2 near root 40 indicate that its left subtree has height 3, right subtree has height 2.
- ◆ By storing the subtree heights an internal node know whether it has become imbalanced





Balanced BST

- ◆ Next we will explain how to perform updates. The most important step is remedy a node u when it becomes imbalanced.
- ◆ It suffices to consider a scenario called 2-level imbalance. In this situation, two conditions apply:
 - ◆ There is a difference of 2 in the heights of the left and right subtree of u .
 - ◆ All the proper descendants of u are balanced
- ◆ We will first explain how to rebalance u in the above situation

2-level imbalance

- There are two cases:


map } red-black tree
BST } splay-tree
AVL-tree




- Due to symmetry, it suffices to explain only the left case, which can be further divide to a left-left and a left-right case, as shown next.

2-level imbalance

- There are two cases:





Left-Left case



Left-Right case


Rebalance Left-Left

- By a rotation:




- Only 3 pointers to change (the red ones). The cost is $O(1)$.
- Recall that $x = h$ or $h+1$

Rebalance Left-Left Example




Rebalance Left-Right

- By a double rotation:



- Only 5 pointers to change (see above). Hence, the cost is $O(1)$.
- Note that x and y must be h or $h-1$. Furthermore, at least one of them must be h (why?)

Rebalance Left-Right Example



Insertion and Deletion Time

◆ Insertion time analysis

- ◆ It will be left as an exercise for you to prove
 - ◆ Only 2-level imbalance can occur in an insertion
 - ◆ Once we have remedied the lowest imbalance node, all the nodes in the tree will become balanced again
- ◆ Thus, we can conclude the insertion cost in a balanced BST is $O(\log n)$, why?

◆ Deletion time analysis

- ◆ It will be left as an exercise for you to prove
 - ◆ Only 2-level imbalance can occur after a deletion
- ◆ Thus, we can conclude the deletion cost in a balanced BST is $O(\log n)$

Balanced BST

- ◆ We now conclude our discussion on the AVL-tree, which provides the following guarantees:
 - ◆ $O(n)$ space consumption
 - ◆ $O(\log n)$ time per predecessor query (hence, also per dictionary lookup)
 - ◆ $O(\log n)$ time per insertion
 - ◆ $O(\log n)$ time per deletion
- ◆ All the above complexities hold in the worst case.

Thank You!